

Introduction

This project aims to predict individuals' likelihood of receiving the H1N1 and seasonal flu vaccines using machine learning models. Using a dataset from DrivenData that includes demographic, behavioral, and health-related features, we frame this as a classification problem. Logistic Regression and Decision Trees are used to build predictive models, with performance evaluated using metrics such as ROC-AUC. The project demonstrates the full data science pipeline, from preprocessing to model interpretation, providing insights into factors influencing vaccine uptake. These findings can inform targeted public health strategies to improve vaccination rates and address vaccine hesitancy.

Business Problem

Despite the availability of vaccines, many individuals choose not to receive them, leaving communities vulnerable to seasonal flu outbreaks and pandemics like H1N1. Public health organizations need to understand the factors influencing vaccination uptake to design effective, targeted interventions.

Key Business Questions:

1. What factors influence an individual's likelihood to receive the H1N1 and seasonal flu vaccines?
2. How can predictive modeling help identify high-risk groups who are less likely to vaccinate?
3. What actionable strategies can public health organizations implement to increase vaccination rates?

Objectives

To develop a predictive model to estimate the likelihood of individuals receiving two specific vaccines: the H1N1 vaccine and the seasonal flu vaccine.

Data

The data for this competition comes from the National 2009 H1N1 Flu Survey (NHFS).

<https://www.drivendata.org/competitions/66/flu-shot-learning/data/>

The source dataset comes with the following data use restrictions:

1. Use the data in these data files for statistical reporting and analysis only.
2. Make no use of the identity of any person or establishment discovered inadvertently and advise the Director, NCHS, of any such discovery (1 (800) 232-4636).
3. Not link these data files with individually identifiable data from other NCHS or non-NCHS data files.

Exploratory Data Analysis

```
# Import the necessary packages
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import roc_auc_score
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score, ConfusionMatrixDisplay
from sklearn.metrics import roc_curve, auc
```

Load the Data

```
# Load the data
train_data=pd.read_csv('D:\\MORINGA\\Phase 3\\Phase 3 Project\\Data\\
training_set_features.csv')
train_labels = pd.read_csv('D:\\MORINGA\\Phase 3\\Phase 3 Project\\
Data\\training_set_labels.csv')
test_data=pd.read_csv('D:\\MORINGA\\Phase 3\\Phase 3 Project\\Data\\
test_set_features.csv')
```

Basic info

```
train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 26707 entries, 0 to 26706
```

```
Data columns (total 36 columns):
```

#	Column	Non-Null Count	Dtype
0	respondent_id	26707 non-null	int64
1	h1n1_concern	26615 non-null	float64
2	h1n1_knowledge	26591 non-null	float64
3	behavioral_antiviral_meds	26636 non-null	float64
4	behavioral_avoidance	26499 non-null	float64
5	behavioral_face_mask	26688 non-null	float64
6	behavioral_wash_hands	26665 non-null	float64
7	behavioral_large_gatherings	26620 non-null	float64
8	behavioral_outside_home	26625 non-null	float64
9	behavioral_touch_face	26579 non-null	float64
10	doctor_recc_h1n1	24547 non-null	float64
11	doctor_recc_seasonal	24547 non-null	float64
12	chronic_med_condition	25736 non-null	float64
13	child_under_6_months	25887 non-null	float64
14	health_worker	25903 non-null	float64
15	health_insurance	14433 non-null	float64
16	opinion_h1n1_vacc_effective	26316 non-null	float64

```

17 opinion_hlnl_risk 26319 non-null float64
18 opinion_hlnl_sick_from_vacc 26312 non-null float64
19 opinion_seas_vacc_effective 26245 non-null float64
20 opinion_seas_risk 26193 non-null float64
21 opinion_seas_sick_from_vacc 26170 non-null float64
22 age_group 26707 non-null object
23 education 25300 non-null object
24 race 26707 non-null object
25 sex 26707 non-null object
26 income_poverty 22284 non-null object
27 marital_status 25299 non-null object
28 rent_or_own 24665 non-null object
29 employment_status 25244 non-null object
30 hhs_geo_region 26707 non-null object
31 census_msa 26707 non-null object
32 household_adults 26458 non-null float64
33 household_children 26458 non-null float64
34 employment_industry 13377 non-null object
35 employment_occupation 13237 non-null object
dtypes: float64(23), int64(1), object(12)
memory usage: 7.3+ MB

```

```
# test_data info
```

```
test_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 26708 entries, 0 to 26707
```

```
Data columns (total 36 columns):
```

#	Column	Non-Null Count	Dtype
0	respondent_id	26708 non-null	int64
1	hlnl_concern	26623 non-null	float64
2	hlnl_knowledge	26586 non-null	float64
3	behavioral_antiviral_meds	26629 non-null	float64
4	behavioral_avoidance	26495 non-null	float64
5	behavioral_face_mask	26689 non-null	float64
6	behavioral_wash_hands	26668 non-null	float64
7	behavioral_large_gatherings	26636 non-null	float64
8	behavioral_outside_home	26626 non-null	float64
9	behavioral_touch_face	26580 non-null	float64
10	doctor_recc_hlnl	24548 non-null	float64
11	doctor_recc_seasonal	24548 non-null	float64
12	chronic_med_condition	25776 non-null	float64
13	child_under_6_months	25895 non-null	float64
14	health_worker	25919 non-null	float64
15	health_insurance	14480 non-null	float64
16	opinion_hlnl_vacc_effective	26310 non-null	float64
17	opinion_hlnl_risk	26328 non-null	float64
18	opinion_hlnl_sick_from_vacc	26333 non-null	float64
19	opinion_seas_vacc_effective	26256 non-null	float64

20	opinion_seas_risk	26209	non-null	float64
21	opinion_seas_sick_from_vacc	26187	non-null	float64
22	age_group	26708	non-null	object
23	education	25301	non-null	object
24	race	26708	non-null	object
25	sex	26708	non-null	object
26	income_poverty	22211	non-null	object
27	marital_status	25266	non-null	object
28	rent_or_own	24672	non-null	object
29	employment_status	25237	non-null	object
30	hhs_geo_region	26708	non-null	object
31	census_msa	26708	non-null	object
32	household_adults	26483	non-null	float64
33	household_children	26483	non-null	float64
34	employment_industry	13433	non-null	object
35	employment_occupation	13282	non-null	object

dtypes: float64(23), int64(1), object(12)
memory usage: 7.3+ MB

Check missing values in train_features

train_data.isnull().sum().sort_values()

respondent_id	0
sex	0
hhs_geo_region	0
census_msa	0
race	0
age_group	0
behavioral_face_mask	19
behavioral_wash_hands	42
behavioral_antiviral_meds	71
behavioral_outside_home	82
behavioral_large_gatherings	87
h1n1_concern	92
h1n1_knowledge	116
behavioral_touch_face	128
behavioral_avoidance	208
household_children	249
household_adults	249
opinion_h1n1_risk	388
opinion_h1n1_vacc_effective	391
opinion_h1n1_sick_from_vacc	395
opinion_seas_vacc_effective	462
opinion_seas_risk	514
opinion_seas_sick_from_vacc	537
health_worker	804
child_under_6_months	820
chronic_med_condition	971
education	1407
marital_status	1408

employment_status	1463
rent_or_own	2042
doctor_recc_h1n1	2160
doctor_recc_seasonal	2160
income_poverty	4423
health_insurance	12274
employment_industry	13330
employment_occupation	13470
dtype: int64	

Dealing with missing values

```
# Fill missing categorical features with the mode
categorical_columns =
train_data.select_dtypes(include=['object']).columns
for col in categorical_columns:
    train_data[col].fillna(train_data[col].mode()[0], inplace=True)
    test_data[col].fillna(train_data[col].mode()[0], inplace=True)

# Fill missing numerical features with the median
numerical_columns = train_data.select_dtypes(include=['float64',
'int64']).columns
for col in numerical_columns:
    train_data[col].fillna(train_data[col].median(), inplace=True)
    test_data[col].fillna(train_data[col].median(), inplace=True)
```

Check for duplicates


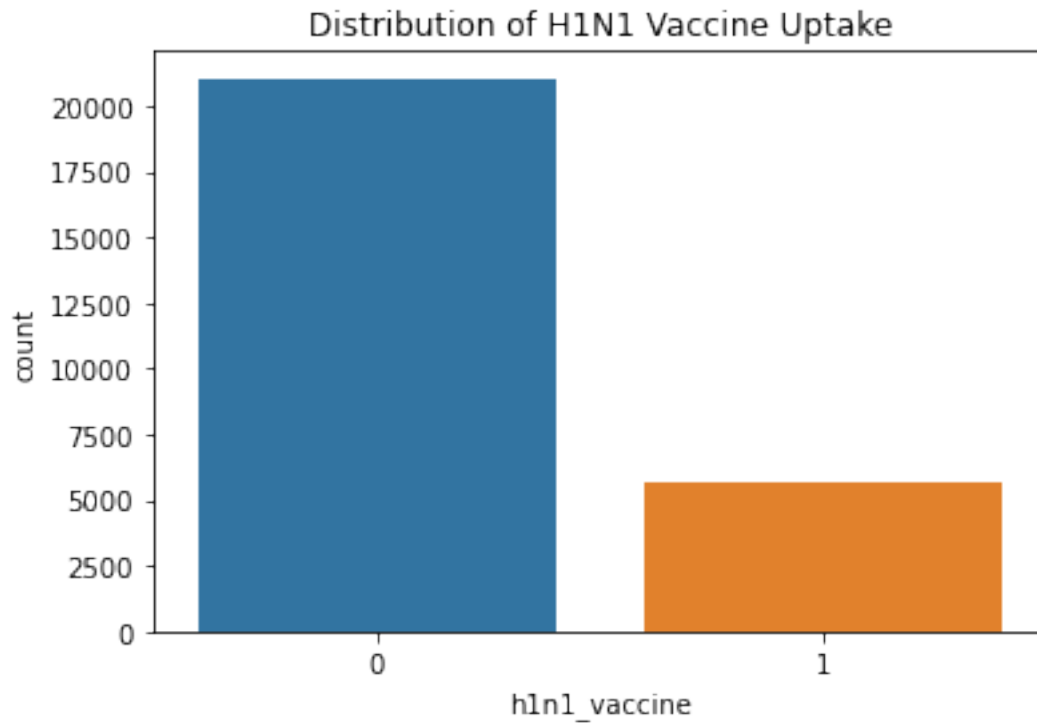
```
train_data.duplicated().sum()

0
```

Analyze target variables

```
# Distribution of vaccine targets
sns.countplot(x='h1n1_vaccine', data=train_labels)
plt.title('Distribution of H1N1 Vaccine Uptake')
plt.show()

sns.countplot(x='seasonal_vaccine', data=train_labels)
plt.title('Distribution of Seasonal Vaccine Uptake')
plt.show()
```



Observation

Most of the respondents didn't get either of the vaccines

Analyze the features

```

# Combine train data and labels
data_combined = pd.concat([train_data, train_labels], axis=1)

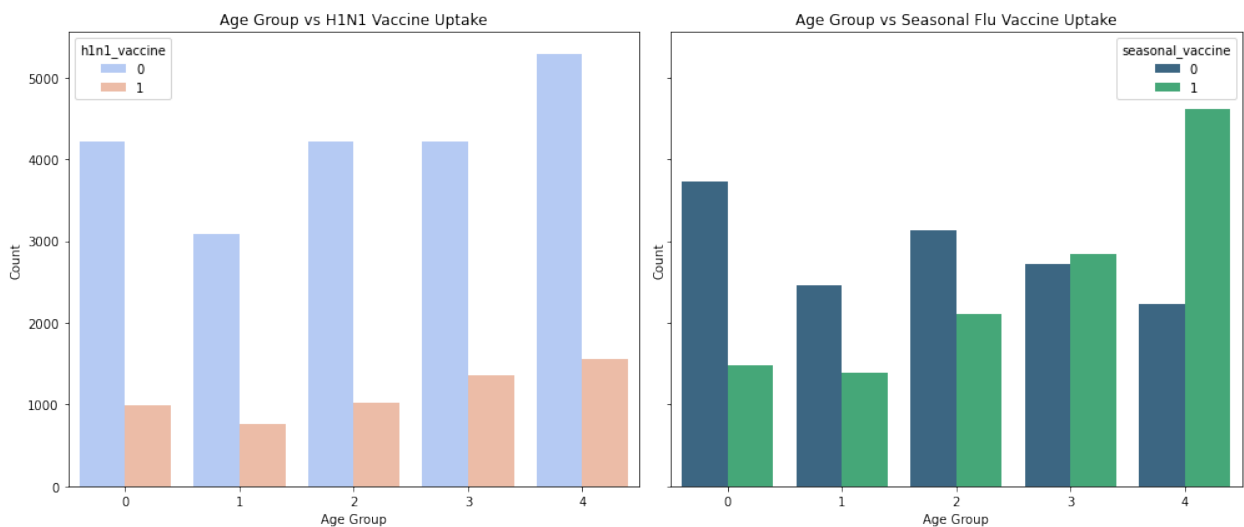
# Set up the figure and axes
fig, axes = plt.subplots(1, 2, figsize=(14, 6), sharey=True)

# H1N1 vaccine uptake
sns.countplot(x="age_group", hue="h1n1_vaccine", data=data_combined,
ax=axes[0], palette="coolwarm")
axes[0].set_title("Age Group vs H1N1 Vaccine Uptake")
axes[0].set_xlabel("Age Group")
axes[0].set_ylabel("Count")

# Seasonal flu vaccine uptake
sns.countplot(x="age_group", hue="seasonal_vaccine",
data=data_combined, ax=axes[1], palette="viridis")
axes[1].set_title("Age Group vs Seasonal Flu Vaccine Uptake")
axes[1].set_xlabel("Age Group")
axes[1].set_ylabel("Count")

# Adjust layout
plt.tight_layout()
plt.show()

```



Observation

A smaller proportion of individuals chose to vaccinate against H1N1. Individuals aged 65 and above had a lower likelihood of receiving the H1N1 vaccine and a higher likelihood of receiving the seasonal vaccine.


```

# Correlation heatmap
correlation = train_data.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation, annot=True, cmap='coolwarm')

```



```
plt.title('Feature Correlation Heatmap')
plt.show()
```



Data Preprocessing

```
# Label encoding for categorical features
encoder = LabelEncoder()
for col in categorical_columns:
    train_data[col] = encoder.fit_transform(train_data[col])
    test_data[col] = encoder.transform(test_data[col])

# Combine features and labels
data = train_data.copy()
data['h1n1_vaccine'] = train_labels['h1n1_vaccine']
data['seasonal_vaccine'] = train_labels['seasonal_vaccine']
```



```
# Split data for H1N1 and seasonal vaccine prediction
X_h1n1 = data.drop(columns=['h1n1_vaccine', 'seasonal_vaccine'])
y_h1n1 = data['h1n1_vaccine']
X_seasonal = data.drop(columns=['h1n1_vaccine', 'seasonal_vaccine'])
y_seasonal = data['seasonal_vaccine']

X_h1n1_train, X_h1n1_val, y_h1n1_train, y_h1n1_val =
train_test_split(X_h1n1, y_h1n1, test_size=0.2, random_state=42)
X_seasonal_train, X_seasonal_val, y_seasonal_train, y_seasonal_val =
train_test_split(X_seasonal, y_seasonal, test_size=0.2,
random_state=42)
```

Model Building

Baseline Model (Logistic Regression Model)

```
# H1N1 Logistic Regression
lr_h1n1 = LogisticRegression(max_iter=1000)
lr_h1n1.fit(X_h1n1_train, y_h1n1_train)
y_h1n1_pred = lr_h1n1.predict(X_h1n1_val)
print("H1N1 Vaccine Logistic Regression Report:")
print(classification_report(y_h1n1_val, y_h1n1_pred))

# Seasonal Logistic Regression
lr_seasonal = LogisticRegression(max_iter=1000)
lr_seasonal.fit(X_seasonal_train, y_seasonal_train)
y_seasonal_pred = lr_seasonal.predict(X_seasonal_val)
print("Seasonal Vaccine Logistic Regression Report:")
print(classification_report(y_seasonal_val, y_seasonal_pred))
```

H1N1 Vaccine Logistic Regression Report:

	precision	recall	f1-score	support
0	0.84	0.95	0.89	4212
1	0.63	0.33	0.43	1130
accuracy			0.82	5342
macro avg	0.73	0.64	0.66	5342
weighted avg	0.80	0.82	0.79	5342

Seasonal Vaccine Logistic Regression Report:

	precision	recall	f1-score	support
0	0.77	0.77	0.77	2891
1	0.73	0.72	0.73	2451
accuracy			0.75	5342
macro avg	0.75	0.75	0.75	5342

weighted avg	0.75	0.75	0.75	5342
--------------	------	------	------	------

Observation

The H1N1 vaccine model has higher accuracy (82%) than the seasonal vaccine model (75%), but its performance for minority Class 1 is suboptimal, suggesting potential room for improvement.

1. For the H1N1 vaccine, the model performs well in identifying individuals who did not receive the vaccine (Class 0), but struggles with accurately predicting those who received it (Class 1), as shown by the lower recall and F1-score for Class 1.
2. For the seasonal vaccine, the model achieves more balanced performance across both classes, with similar precision, recall, and F1-scores.

Decision Tree Classifier

```
# Initialize Decision Tree classifiers
dt_h1n1 = DecisionTreeClassifier(random_state=42)
dt_seasonal = DecisionTreeClassifier(random_state=42)

# Train the models
dt_h1n1.fit(X_h1n1_train, y_h1n1_train)
dt_seasonal.fit(X_seasonal_train, y_seasonal_train)

# Predict probabilities
dt_h1n1_probs = dt_h1n1.predict_proba(X_h1n1_val)[: , 1]
dt_seasonal_probs = dt_seasonal.predict_proba(X_seasonal_val)[: , 1]

# Calculate AUC
h1n1_auc = roc_auc_score(y_h1n1_val, dt_h1n1_probs)
seasonal_auc = roc_auc_score(y_seasonal_val, dt_seasonal_probs)
print(f"H1N1 Vaccine AUC: {h1n1_auc}")
print(f"Seasonal Vaccine AUC: {seasonal_auc}")
```

```
H1N1 Vaccine AUC: 0.6365241745035256
Seasonal Vaccine AUC: 0.676569090387436
```

Observation

The Seasonal vaccine model is performing slightly better than the H1N1 vaccine model based on these AUC scores. However, both are somewhat lower than ideal, indicating there may be room for improvement in the models' predictive accuracy.

Model Evaluation

```
# Train logistic regression for H1N1 vaccine
logreg_h1n1 = LogisticRegression(max_iter=1000, random_state=42)
logreg_h1n1.fit(X_h1n1_train, y_h1n1_train)
```

```

# Train logistic regression for seasonal vaccine
logreg_seasonal = LogisticRegression(max_iter=1000, random_state=42)
logreg_seasonal.fit(X_seasonal_train, y_seasonal_train)

# Generate ROC curves for H1N1 vaccine
fpr_h1n1, tpr_h1n1, _ = roc_curve(y_h1n1_val,
logreg_h1n1.predict_proba(X_h1n1_val)[: , 1])
roc_auc_h1n1 = auc(fpr_h1n1, tpr_h1n1)

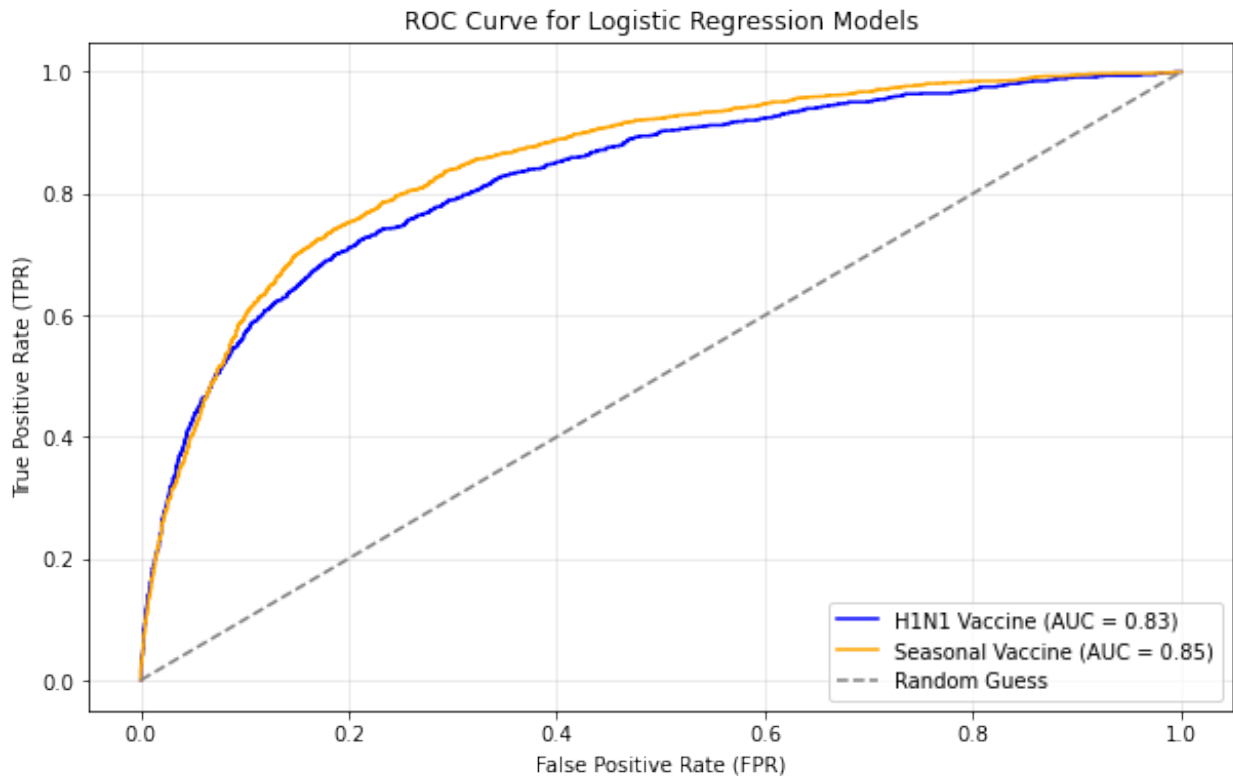
# Generate ROC curves for seasonal vaccine
fpr_seasonal, tpr_seasonal, _ = roc_curve(y_seasonal_val,
logreg_seasonal.predict_proba(X_seasonal_val)[: , 1])
roc_auc_seasonal = auc(fpr_seasonal, tpr_seasonal)

# Plot the ROC curves
plt.figure(figsize=(10, 6))
plt.plot(fpr_h1n1, tpr_h1n1, label=f"H1N1 Vaccine (AUC =
{roc_auc_h1n1:.2f})", color='blue')
plt.plot(fpr_seasonal, tpr_seasonal, label=f"Seasonal Vaccine (AUC =
{roc_auc_seasonal:.2f})", color='orange')

# Plot the diagonal line for random predictions
plt.plot([0, 1], [0, 1], color="grey", linestyle="--", label="Random
Guess")

# Customize the plot
plt.title("ROC Curve for Logistic Regression Models")
plt.xlabel("False Positive Rate (FPR)")
plt.ylabel("True Positive Rate (TPR)")
plt.legend(loc="lower right")
plt.grid(alpha=0.3)
plt.show()

```




```
# Plot ROC curves
plt.figure(figsize=(12, 6))

# H1N1 ROC Curve
fpr_h1n1, tpr_h1n1, _ = roc_curve(y_h1n1_val, dt_h1n1_probs)
plt.plot(fpr_h1n1, tpr_h1n1, label=f"H1N1 Vaccine (AUC = {h1n1_auc:.2f})")

# Seasonal ROC Curve
fpr_seasonal, tpr_seasonal, _ = roc_curve(y_seasonal_val, dt_seasonal_probs)
plt.plot(fpr_seasonal, tpr_seasonal, label=f"Seasonal Vaccine (AUC = {seasonal_auc:.2f})")

# Plot settings
plt.plot([0, 1], [0, 1], 'k--', label='Random Chance')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves for Decision Tree Classifiers')
plt.legend()
plt.show()
```



Observation

Both models have AUCs greater than 0.5, indicating that both are performing better than random guessing. However, neither model has an exceptionally high AUC, suggesting there is room for improvement. The Seasonal Vaccine model is performing slightly better than the H1N1 Vaccine model, as its curve is higher, and its AUC is 0.68 compared to 0.64 for H1N1.

Compare the Models

```
from sklearn.metrics import accuracy_score,
precision_recall_fscore_support, roc_auc_score, confusion_matrix,
log_loss

# Logistic Regression for H1N1 Vaccine
h1n1_lr_auc = roc_auc_score(y_h1n1_val,
lr_h1n1.predict_proba(X_h1n1_val)[: , 1])
h1n1_lr_fpr, h1n1_lr_tpr, _ = roc_curve(y_h1n1_val,
lr_h1n1.predict_proba(X_h1n1_val)[: , 1])

# Decision Tree for H1N1 Vaccine
h1n1_dt_auc = roc_auc_score(y_h1n1_val,
dt_h1n1.predict_proba(X_h1n1_val)[: , 1])
h1n1_dt_fpr, h1n1_dt_tpr, _ = roc_curve(y_h1n1_val,
dt_h1n1.predict_proba(X_h1n1_val)[: , 1])

# Logistic Regression for Seasonal Vaccine
seasonal_lr_auc = roc_auc_score(y_seasonal_val,
lr_seasonal.predict_proba(X_seasonal_val)[: , 1])
seasonal_lr_fpr, seasonal_lr_tpr, _ = roc_curve(y_seasonal_val,
```

```

lr_seasonal.predict_proba(X_seasonal_val)[:, 1])

# Decision Tree for Seasonal Vaccine
seasonal_dt_auc = roc_auc_score(y_seasonal_val,
dt_seasonal.predict_proba(X_seasonal_val)[:, 1])
seasonal_dt_fpr, seasonal_dt_tpr, _ = roc_curve(y_seasonal_val,
dt_seasonal.predict_proba(X_seasonal_val)[:, 1])

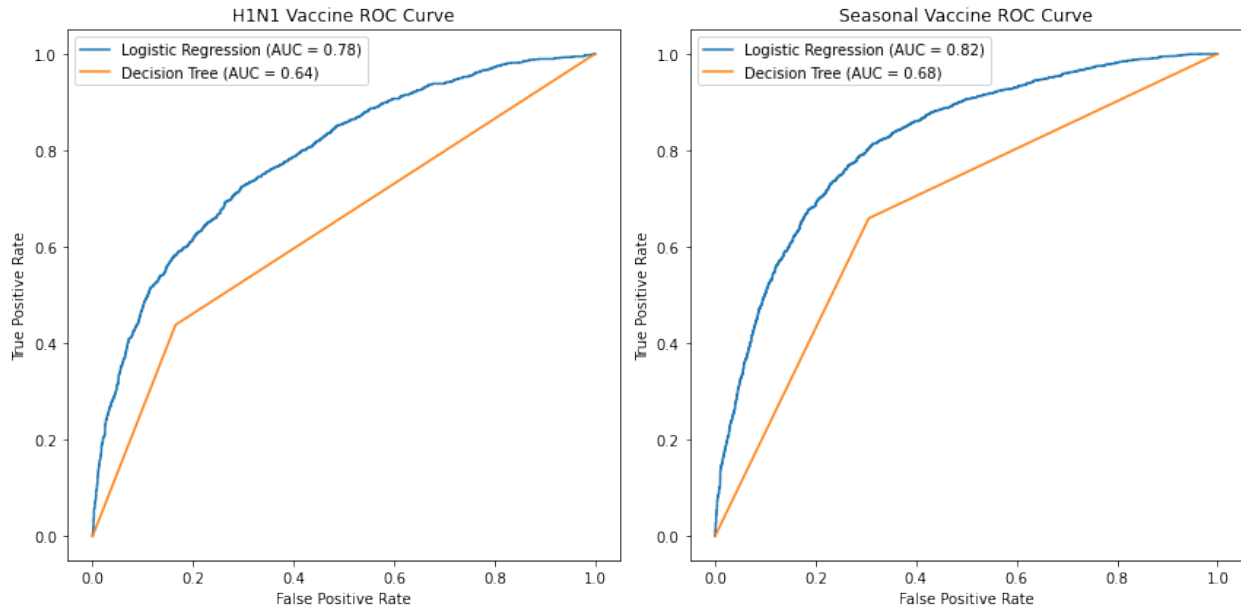
# Plot ROC Curves
plt.figure(figsize=(12, 6))

# H1N1 Vaccine
plt.subplot(1, 2, 1)
plt.plot(h1n1_lr_fpr, h1n1_lr_tpr, label=f"Logistic Regression (AUC =
{h1n1_lr_auc:.2f})")
plt.plot(h1n1_dt_fpr, h1n1_dt_tpr, label=f"Decision Tree (AUC =
{h1n1_dt_auc:.2f})")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("H1N1 Vaccine ROC Curve")
plt.legend()

# Seasonal Vaccine
plt.subplot(1, 2, 2)
plt.plot(seasonal_lr_fpr, seasonal_lr_tpr, label=f"Logistic Regression
(AUC = {seasonal_lr_auc:.2f})")
plt.plot(seasonal_dt_fpr, seasonal_dt_tpr, label=f"Decision Tree (AUC
= {seasonal_dt_auc:.2f})")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Seasonal Vaccine ROC Curve")
plt.legend()

plt.tight_layout()
plt.show()

```



Observation

logistic Regression consistently outperforms Decision Tree in both tasks, although the difference in AUC is not significant. Both models show good performance, with AUC values above 0.75 for both vaccines, which is considered acceptable for classification tasks.

Feature Importance

```
# Drop the 'respondent_id' column from the training data
X_h1n1_train = X_h1n1_train.drop(columns=['respondent_id'])
X_h1n1_val = X_h1n1_val.drop(columns=['respondent_id'])
X_seasonal_train = X_seasonal_train.drop(columns=['respondent_id'])
X_seasonal_val = X_seasonal_val.drop(columns=['respondent_id'])

# Train the Decision Tree Classifier again
dt_h1n1 = DecisionTreeClassifier()
dt_h1n1.fit(X_h1n1_train, y_h1n1_train)

# Get feature importances
feature_importance = dt_h1n1.feature_importances_

# Create a DataFrame for better visualization
feature_names = X_h1n1_train.columns
importance_df = pd.DataFrame({
    'Feature': feature_names,
    'Importance': feature_importance
})

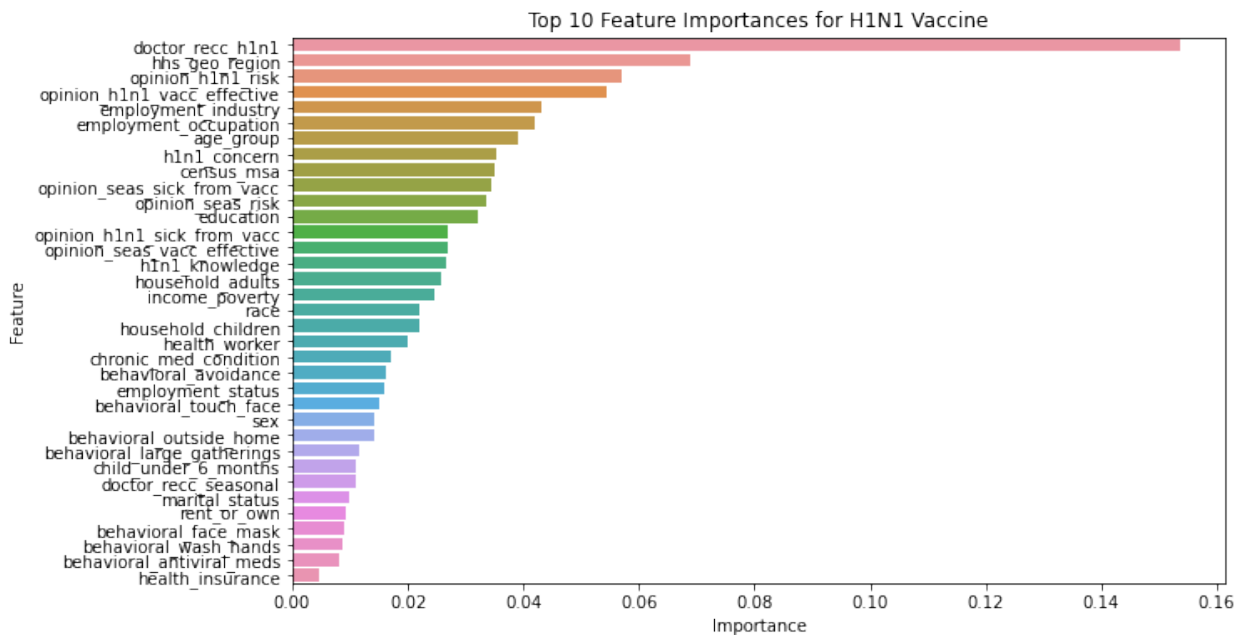
# Sort the features by importance
importance_df = importance_df.sort_values(by='Importance',
ascending=False)
```



```
# Display the top 10 most important features
print(importance_df.head(10))

# Plot the feature importances
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=importance_df)
plt.title('Top 10 Feature Importances for H1N1 Vaccine')
plt.show()
```

	Feature	Importance
9	doctor_recc_h1n1	0.153767
29	hhs_geo_region	0.068879
16	opinion_h1n1_risk	0.056995
15	opinion_h1n1_vacc_effective	0.054501
33	employment_industry	0.043200
34	employment_occupation	0.041958
21	age_group	0.039056
0	h1n1_concern	0.035230
30	census_msa	0.034950
20	opinion_seas_sick_from_vacc	0.034436



Observation

The feature importance values indicate that the most significant factors influencing H1N1 vaccination are:

1. Doctor recommendation (doctor_recc_h1n1): Strongly influences the likelihood of vaccination, suggesting that healthcare professional advice is crucial.

2. Geographic region (hhs_geo_region): Regional variations likely affect vaccination rates, potentially due to different healthcare access or public health initiatives.
3. Opinions on risk and vaccine effectiveness: These factors suggest that individuals' perceptions of risk and the effectiveness of the vaccine are significant in their vaccination decision.

Conclusions

1. Factors influencing vaccination Demographic factors such as age, significantly affect vaccine uptake. For instance, older groups were more likely to skip vaccination. Behavioral and informational aspects, like trust in healthcare providers and access to reliable information, also play key roles.
2. Model Performance Both logistic regression and decision tree models provided reasonable performance, with AUC scores above 0.5 for both H1N1 and seasonal vaccines. Logistic regression models excelled in interpretability, making them suitable for identifying key predictors. Decision trees offered a transparent, rule-based approach but showed lower AUCs, suggesting that they might require further tuning or ensemble methods for optimal performance.
3. Strategies for Stakeholders Public health organizations should focus on tailored interventions, such as: Educational campaigns aimed at dispelling vaccine misinformation. Community outreach programs to improve access to vaccines for underserved populations. Healthcare provider training to communicate vaccine benefits effectively.

Recommendations

1. Targeted interventions: Focus on educating people about vaccine effectiveness and safety, especially in regions with lower vaccination rates.
2. Leverage healthcare providers: Strengthen communication between healthcare professionals and individuals to boost vaccine uptake.